


F1 per model when the judge reads only the embedded text vs when it also sees the image.

F1 · what the image adds (prompt v1) [95 % CI]



Error bars: bootstrap 95 % CIs. 50 of 60 positives contain an explicit keyword: the benchmark mostly measures reading the embedded text.