


F1 and AUROC per model with bootstrap 95 % CIs: the four models are within each other's intervals.

Four vision models on sv_images_v1 (60 / 60, prompt v1 frozen, image + text)



Error bars: bootstrap 95 % CIs (2 000 resamples). Differences are not statistically significant (CIs overlap, McNemar $p \geq 0.6$).