



Qwen3.5-9B · prompt v1 vs v3



F1 · what the image adds (prompt v1)



v3 was tuned after analysing this set's errors: indicative only. 50 of 60 positives contain an explicit keyword: the benchmark mostly measures reading the embedded text.