



Left: the tuned prompt v3 against the frozen v1; right: F1 with the embedded text alone vs with the image added, per model.

Qwen3.5-9B · prompt v1 vs v3



F1 · what the image adds (prompt v1) [95 % CI]



Error bars: bootstrap 95 % CIs; no error bar = CI not available (IC non disponible) (prompt v3; AUROC text only). v3 was tuned after analysing this set's errors: indicative only. 50 of 60 positives contain an explicit keyword: the benchmark mostly measures reading the embedded text.