


sv_images_v1 (60 sexual / 60 not sexual), prompt v1 frozen, image + text: the four models score within each other's confidence intervals.

Four vision models · F1 and AUROC [95 % CI]



Error bars: bootstrap 95 % CIs (2 000 resamples of the 120 images). McNemar vs Qwen3.5-9B: $p \geq 0.6$ for all $\rightarrow$ no significant difference.