


All four AI models' confidence separates harmful from harmless memes very well

Four open-source AI models · 120 hand-checked memes



Same 120 hand-checked memes. Besides its yes / no answer, each model gives a confidence score. Bars show the ranking quality (AUROC): how well this confidence puts harmful memes above harmless ones (1.0 = perfect, 0.5 = chance). Black lines show the uncertainty range.