

InternVL3.5-8B catches 57 of 60 harmful memes and raises 10 false alarms

Four open-source AI models · 120 hand-checked memes

		What the AI decided	
What a human decided		Flagged	Not flagged
Harmful (60)	Harmful (60)	57 Correctly flagged	3 Missed
		10 False alarm	50 Correctly ignored

We showed 120 memes, checked by hand (60 with sexual content, 60 without), to each AI model, which saw the image and the text written on it. Rows: what a human decided. Columns: what the AI decided. Green: right answers; orange: mistakes. Our definition, first version, frozen before the test.