


The AI's confidence separates racist from non-racist videos perfectly

Qwen3.5-9B · 30 public videos (15 racist, 15 not)



Each video gets a confidence score (its most suspicious scene). Lowering the alert threshold moves along the curve: up = more racist videos caught, right = more harmless videos flagged. A curve through the top-left corner is perfect (ranking quality 1.0); the dashed line is chance. Both setups give the same curve. The video test uses racism because no open video dataset on sexual violence exists.