

A fast pre-filter halves the work of the AI judge, missing 2 of 60 harmful memes

Simulated feed · 579 memes · share of harmful content in the feed: 10 %

What a human decided		What the pre-filter decided	
		Passed to the AI judge	Set aside
Harmful (60)		58 Passed on	2 Missed
Harmless (519)		219 Extra work for the judge	300 Correctly set aside

Two-step system: a fast pre-filter looks at each image first and passes only suspicious memes to the AI judge (Qwen3.5-9B). Feed: the 120 hand-checked memes plus 459 harmless ones, so 10 % is harmful. The pre-filter passed 48 % of the memes on. The 459 extra memes were labelled by the AI judge, not checked by hand: this is an estimate.