

When half of the memes are harmful, the pre-filter saves only about a quarter of the work

120 hand-checked memes · share of harmful content: 50 %

		What the pre-filter decided	
What a human decided		Passed to the AI judge	Set aside
Harmful (60)	Harmful (60)	58 Passed on	2 Missed
		30 Extra work for the judge	30 Correctly set aside

Two-step system: a fast pre-filter looks at each image first and passes only suspicious memes to the AI judge (Qwen3.5-9B). On the 120 hand-checked memes (half harmful), it passed 73 % of them on and missed 2 harmful memes. Real feeds contain far less harmful content, so the saving is larger there (see the 10 % feed).