


The pre-filter learns much better from the AI judge than from the original dataset labels

Three fast pre-filters · tested on the 120 hand-checked memes



Each pre-filter learned from other memes (never these 120), labelled either by the original dataset (grey) or by the AI judge (yellow). Bars show the ranking quality (AUROC): how well its score puts harmful memes above harmless ones (1.0 = perfect, 0.5 = chance). No uncertainty range available.